

An Analysis of Parallel Rendering Systems

Stefan Eilemann*

This White Paper analyzes and classifies the different approaches taken to parallel rendering. The purpose is to clarify common misconceptions and false expectations of the capabilities of the different classes of parallel rendering software.

We examine the rendering pipeline of the typical visualization application and identify the typical bottlenecks found in this pipeline. The methods used for parallel rendering are classified in three fundamental approaches and then analyzed with respect to their influence on this rendering pipeline, which leads to conclusions about possible performance gains and the necessary porting effort.

We advocate the need for a generic, open parallel rendering framework to build scalable graphics software. The inherent limitations of transparent multi-pipe software are justified, and a full software stack for graphics clusters is outlined.

Version	Date	Changes
0.5	May 19, 2006	Initial Version

*eilemann@gmail.com

1. Introduction

The purpose of this white paper is to argue for the need for a middleware to build truly scalable, interactive visualization applications. In contrast to high-performance computing (HPC), the high-performance visualization (HPV) community has not generally accepted the need to parallelize applications for performance. The belief that scalability can be achieved transparently by some magic software is still wide-spread, and contributes to the slow adoption of graphics clusters and virtual reality installations due to the application performance and (un-)availability.

1.1. Motivation

Application developers are faced with making their applications aware of multiple graphics cards for two reasons – either the application has to run on advanced display systems, or the rendering has to be accelerated by using parallel rendering on multiple graphics cards.

1.1.1. Multi-View Rendering

The first motivation is often to run the application on advanced display systems: workstations with multiple monitors, walls build out of multiple displays or projectors as well as immersive environments. Such installations need more than one graphics card to drive all the display systems. Therefore the rendering has to be distributed across a number of cards and often across multiple systems. The outputs of the displays has to be synchronized spatially and time-wise to provide one coherent image. In immersive environments, stereo rendering is a requirement.




Figure 1: Some advanced display systems

Applications which scale their display size often have to scale their rendering performance, because the increased display size exposes previously undiscovered bottlenecks in the rendering pipeline. The higher display resolution increases the needed pixel fill rate. Due to the increased image fidelity, data of a higher resolution is used – which increases the requirements on the whole rendering pipeline.

1.1.2. Scalable Rendering

The second motivation is to accelerate the rendering by aggregating multiple resources, because the application hits limitations of a single graphics system.

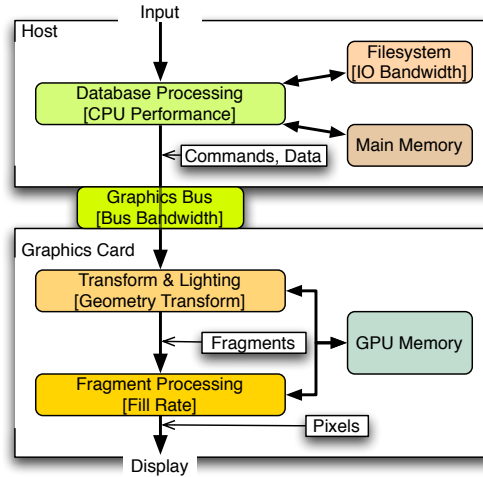


Figure 2: The basic rendering pipeline

Figure 2 shows a simplified rendering pipeline and the resources used by during rendering. The typical performance bottlenecks in such a pipeline are:

Fill Rate The amount of rendered pixels can exceed the processing capabilities of the graphics card's rasterization pipeline, that is, the amount of pixels the card can render in a given time. Volume rendering and complex fragment shaders easily limit the application's performance due to a fill rate bottleneck.

Geometry Transform Similar to the fill rate bottleneck, but less common, are bottlenecks in the GPU's vertex processing capabilities, that is, the amount of triangles a card can process. Complex geometry and procedural computations in the vertex shader can cause this bottleneck.

GPU Memory Certain applications, for example volume rendering and terrain visualization, easily exceed the amount of memory available on the graphics card. For example, a volume of size 1024^3 has a size of 1 GB if it is black and white, and 4 GB if it is colored.

Bus Bandwidth Dynamic data, such as time-varying data sets and interactive modifications of the model, has high requirements on the system's IO capabilities. The data can not be cached on the GPU's memory but has to be transferred to the card from main memory (or disk) during rendering. This is a common bottleneck of interactive applications and leads to a rendering performance below the GPU's capabilities.

CPU Performance Similar to the bus bandwidth, the CPU can become the bottleneck when traversing the database to generate the rendering commands. Tessellation, visibility computations and other interpretation of the application data for rendering cause this bottleneck.

Main Memory The actual model database often holds additional application-specific information, which is not used for rendering. Additionally, the amount of data displayed is drastically reduced during rendering to allow interactive framerates. This leads to application scenarios where the required main memory size exceeds the capabilities of a single system.

IO Bandwidth Data sets which exceed the main memory size are visualized by roaming through the data, that is, a subset of the whole data set is loaded into main memory to be processed and rendered. As the user roams through the data, new pieces of it are loaded from storage, which can be a bottleneck.

1.1.3. Parallelism in Graphics Hardware

Graphics cards address some bottlenecks through hardware parallelism. GPU's have multiple vertex pipelines (geometry transform) and multiple fragment pipelines (fill rate) which accelerate the fill rate and geometry transform rate in hardware, transparently to the application. Due to physical limitations, the number of vertex and fragment shading pipelines on a single GPU is limited.

Recent implementations of multi-GPU systems (nVidia SLI, ATI CrossFire) provide additional parallelism. The individual graphic processors produce a partial image which is then composited using the hardware support. The individual GPU's can sit on the same card, or on separate cards with a special connection between them.

In such a setup, the graphics library dispatches the rendering commands to all rendering units, a process which is transparent to the application. The overall pixel fill rate is increased, while other potential bottlenecks are not addressed, since all commands are sent to all cards and each card needs to have a copy of all the data (textures, display lists, etc.). A special case is the alternate frame rendering (AFR) mode, where each card renders full, alternating frames. This mode also scales geometry transform and bus bandwidth, while increasing the latency between input and the rendered output.

Since the application is unmodified, all rendering commands are still submitted by one thread, which has to be able to saturate the additional graphics cards to make use of the additional hardware. Distributed scene graphs and parallel rendering frameworks can in theory parallelize the rendering for such hardware, once the hardware vendors provide the necessary API's to program the hardware.

2. Transparent Multi-Pipe Rendering Software

Transparent rendering software replaces the system's rendering library, e.g. OpenGL or Direct3D, with a new 'intercept' library. The application's output is taken at the rendering command level, that is, the entry points of the rendering library are replaced by new functions. The new functions package the rendering command stream, analyze it and send it to a number of rendering processes, potentially over the network to other systems. This approach is typically used for the applications which want to use multiple graphics cards to scale the display size. A limitation of this approach is that unmodified applications may not be able to run on non-planar display systems, because important data is not rendered due to the application's view frustum culling.

Transparent solutions can only increase the rendering performance if the main bottleneck is on the graphics card (geometry transform or fill rate). The processing stages to

produce the rendering commands are untouched and remain single-threaded. Potential CPU bottlenecks in the rendering thread are amplified, since packaging and sending of the rendering commands to other processes is more time-consuming than writing the commands directly to a GPU. The bandwidth and latency to the rendering nodes is worse than the direct local connection to the graphics card. In reality, this approach often scales the display size at the expense of performance, since applications are rarely limited by the GPU processing speed.

Analyzing the command stream allows some optimizations. The transparent library can determine the visibility of rendering commands and textures for each rendering unit, and only send the data to the appropriate processes. This can reduce the rendering load of the individual graphics cards. However, this command stream filtering is only beneficial for the performance if it can be performed faster than the actual rendering – which is normally only the case for static chunks of data, such as display lists and large textures.

Some parallel rendering frameworks support parallelization of the application, that is, the application is modified and renders using multiple threads. This support is often rudimentary and does not address other problems, such as configuration and data distribution. Furthermore, by working on the low-level graphics command stream, the attainable performance is limited up-front. From an application developers point of view, writing a parallel application using a transparent framework is at least the same effort as using a parallel rendering framework. For the purpose of this white paper, such an extension to a transparent software can be considered a parallel rendering framework.

3. Distributed Scene Graphs

The distributed scene graph approach requires that the application uses the given scene graph to describe the model database. The rendering of that scene graph is parallelized by the scene graph’s implementation, to a great extent transparently to the application. The scene graph is replicated and kept up-to-date on all rendering nodes. During rendering, the individual instances are traversed in parallel to generate the rendering commands, which are directly send to the graphics hardware.

From the applications point of view, this approach is similar to the transparent solutions, in that the rendering is parallelized. The big difference is that a distributed scene graph approaches the problem at a much higher level, which allows significant optimisations and leads to better performance. In contrast to transparent rendering software, the traversal to generate the rendering commands is parallelized, and the commands are send –in parallel– directly to the graphics card. The scene graph frees the application developer of the task to parallelize the rendering.

4. Parallel Rendering Frameworks

Parallel rendering frameworks help application developers to parallelize their application. They do not impose a scene graph or other rendering libraries on the application. By solving the common problems of multi-pipe rendering, the application developer

can focus on application-specific problems. Parallel rendering frameworks address synchronization, task decomposition and load-balancing, data transport as well as the composition of the rendering result. In many ways, they are comparable to HPC libraries such as MPI and PVM, while being focused on interactive applications and parallel rendering.

Parallel rendering frameworks provide an execution environment for any application, regardless of the used rendering software. Refactoring the application, other than strictly needed for multi-pipe rendering, is not necessary. The application provides entry points for its rendering functions. Depending on the current configuration, the framework creates the necessary rendering threads which are initialised with a rendering context by the application's init functions. When the application request a new frame of rendering, the framework calls the necessary rendering callbacks in the correct context and order.

5. Conclusion

From a technical point of view, choosing the right approach to parallel rendering depends on the requirements and bottlenecks of the application. Figure 3 gives an overview how well the different bottlenecks can be addressed with the various solutions.

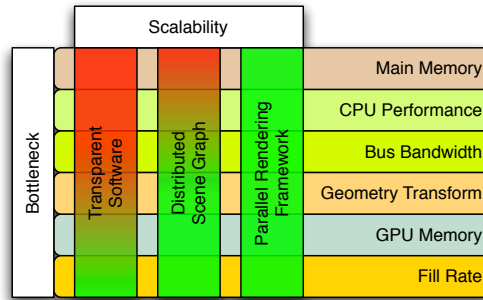


Figure 3: Scalability of the different parallel rendering approaches

The least intrusive approach of using a transparent software is often explored first, which frequently does not meet the expectations and requirements. Performance gains with transparent solutions are possible with benchmarks or well-behaved applications, but real-world software often does not run faster, slowdowns are commonly observed. Transparent rendering software often supports only a subset of the features and extensions of the rendering API, in particular anything which requires a roundtrip to the graphics card is unsupported or slow. On the other hand, a transparent solution is a viable way to run a certain set of applications on graphics clusters. For legacy applications, which cannot be modified, it is the only possibility to make them multi-pipe ready. The low porting effort is a strong benefit for transparent rendering software.

Distributed scene graphs are the ideal and obvious solution if the application already uses such a scene graph, or is planning on using one in the near future. The distributed scene graph, especially when build upon a parallel rendering framework, can deliver

optimal performance while requiring little to no application changes. Porting an existing application to a distributed scene graph is often impractical due to the high porting cost.

Parallel rendering frameworks are the common middleware for any multi-pipe application. Distributed scene graphs and transparent rendering software can use it as the base for their implementation. Applications not suitable for these approaches can easily implement their parallel rendering on such a framework. The porting effort is reduced to the unavoidable refactoring needed to separate the rendering from the core application in order to make it distributable. A parallel rendering framework addresses common problems encountered when doing such a port and follows the natural programming model for any multi-pipe application. For high-performance applications, a parallel rendering framework is the only solution short of a completely custom implementation.

Parallel rendering frameworks can and should be the foundation of transparent solutions and distributed scene graphs. These approaches then become 'applications' of the parallel rendering framework to address a certain subset of parallel rendering applications - namely performance-uncritical and legacy applications, as well as applications already written using a specific scene graph.




Figure 4: Ideal software layering for a visualization system

The three approaches to parallel rendering complement each other and serve different application needs. Transparent solutions are a quick way to run interactive applications on graphic clusters, but they do not provide the necessary performance for many applications. Likewise, distributed scene graphs are often too invasive since they require the applications to use a certain format –the scene graph– to describe their data.

A. Current Parallel Rendering Software

This appendix gives a short and incomplete overview of the existing software solutions for parallel rendering.

A.1. Transparent Multi-Pipe Rendering Software

A.1.1. Chromium

Chromium is a Open Source solution originally developed by the Stanford University. Some graphics cluster vendors, most notably Tungsten Graphics, support and extend Chromium on their hardware platforms. Chromium provides functionality for

the synchronization of the rendering commands when using multiple, parallel application instances.

A.1.2. VGP Software Solutions

The Virtual Graphics Platform sold by ModViz, Inc. is a proprietary software solution similar to Chromium. The VGP Integration API allows to augment the OpenGL command stream with additional information to improve the rendering performance.

A.1.3. OpenGL Multipipe

OpenGL Multipipe is a transparent software solution delivered with SGI's multi-pipe machines. It is not sold separately. For building parallel applications, SGI offered OpenGL Multipipe SDK.

A.2. Distributed Scene Graphs

A.2.1. TGS Open Inventor Cluster Edition

The Open Inventor Cluster Edition is an extension to the Open Inventor scene graph to support parallel, distributed rendering on graphics clusters. It does support both multi-view and scalable rendering.

A.3. Parallel Rendering Frameworks

A.3.1. OpenGL Multipipe SDK

OpenGL Multipipe SDK from Silicon Graphics, Inc. is a proprietary framework for the development of scalable graphics applications. It does not support distributed rendering.

A.3.2. CAVELib

CAVELib from VRCO Inc. is a proprietary, parallel rendering framework which supports parallel multi-view rendering.

A.3.3. Equalizer

Equalizer is an open source parallel rendering framework under development by the University of Zürich and others. Equalizer will support multi-view and scalable rendering, together with a resource management system and a transparent rendering layer.

OpenGL, Open Inventor and OpenGL Multipipe is a trademark of Silicon Graphics, Inc. SLI is a trademark of nVidia. CrossFire is a trademark of ATI. Virtual Graphics Platform is a trademark of ModViz, Inc. CAVELib is a trademark of VRCO Inc. All other products named are trademarks of their respective owners.